

Mini-Project

Black Friday: How much will a customer spend?

Introduction

A retail company “ABC Private Limited” wants to understand the customer purchase behaviour (specifically, purchase amount) against various products of different categories. They have shared purchase summary of various customers for a selected high volume products from last month.

They want to build a model to predict the purchase amount of customer against various products which will help them to create personalized offer for customers against different products.

Performance Measure

Usually for regression problems the typical performance measure is the Root Mean Square Error (RMSE). This function gives an idea of how much error the system makes in its predictions with higher weight for large errors.

Make Assumptions

Before even looking at the available data is good to make some assumption on the expected results. Therefore, let's start to think about possible parameters that might influence the amount a client spends on Black Friday.

Available data

This is the current data available:

```
In [1]: # data.png
```

If we analyse it individually we see that we do not have any information regarding the stores. Moreover, there is some information related to the customer such as age group, sex, occupation and marital status. On the other hand, we have data on the city's size and how many years the customer has lived in it whereas on the product's side there is only information regarding the categories and the amount spent. It is my belief that Gender , Age , City_Category , Product_Category_1 are the predictors that will influence more the amount spent by a customer on this day.

The target variable is Purchase .

```
In [2]: # target.png
```

Take a quick look at the Data Structure

Let's start by importing some libraries and our data.

```
In [3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [4]: train = pd.read_csv("train.csv")
test = pd.read_csv("test-comb.csv")
train.head()
```

Out[4]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_
0	1000001	P00069042	F	0-17	10	A	2
1	1000001	P00248942	F	0-17	10	A	2
2	1000001	P00087842	F	0-17	10	A	2
3	1000001	P00085442	F	0-17	10	A	2
4	1000002	P00285442	M	55+	16	C	4+

In [5]: train.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 12 columns):
User_ID                550068 non-null int64
Product_ID            550068 non-null object
Gender                550068 non-null object
Age                  550068 non-null object
Occupation            550068 non-null int64
City_Category        550068 non-null object
Stay_In_Current_City_Years  550068 non-null object
Marital_Status        550068 non-null int64
Product_Category_1    550068 non-null int64
Product_Category_2    376430 non-null float64
Product_Category_3    166821 non-null float64
Purchase              550068 non-null int64
dtypes: float64(2), int64(5), object(5)
memory usage: 50.4+ MB

```

In [6]: train.describe()

Out[6]:

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Cate
count	5.500680e+05	550068.000000	550068.000000	550068.000000	376430.000000
mean	1.003029e+06	8.076707	0.409653	5.404270	9.842329
std	1.727592e+03	6.522660	0.491770	3.936211	5.086590
min	1.000001e+06	0.000000	0.000000	1.000000	2.000000
25%	1.001516e+06	2.000000	0.000000	1.000000	5.000000
50%	1.003077e+06	7.000000	0.000000	5.000000	9.000000
75%	1.004478e+06	14.000000	1.000000	8.000000	15.000000
max	1.006040e+06	20.000000	1.000000	20.000000	18.000000



Organisation of our analysis

Our goal as a Data Scientist is to identify the most important variables and to define the best regression model for predicting out target variable. Hence, this analysis will be divided into five stages:

1. Exploratory data analysis (EDA);
2. Data Pre-processing;
3. Feature engineering;
4. Feature Transformation;
5. Modeling;

The following is a workflow chart illustrating the five stages:

In [7]: `#workflow.png`

1. Exploratory Data Analysis (EDA)

We've made our first assumptions on the data and now we are ready to perform some basic data exploration and come up with some inference. Hence, the goal for this section is to take a glimpse on the data as well as any irregularities so that we can correct on the next section, Data Pre-Processing.

1.1. Univariate Analysis

To get an idea of the distribution of numerical variables, histograms are an excellent starting point. Let's begin by generating one for Purchase, our target variable.


1.1.1. Distribution of the target variable: Purchase

```
In [8]: plt.style.use('fivethirtyeight')
plt.figure(figsize=(12,7))
sns.distplot(train.Purchase, bins = 25)
plt.xlabel("Amount spent in Purchase")
plt.ylabel("Number of Buyers")
plt.title("Purchase amount Distribution")
```

C:\Users\Mayank\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: Future Warning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[8]: Text(0.5, 1.0, 'Purchase amount Distribution')
```



```
In [9]: print ("Skew is:", train.Purchase.skew())
print("Kurtosis: %f" % train.Purchase.kurt())
```

```
Skew is: 0.6001400037087128
```

```
Kurtosis: -0.338378
```

1.1.2. Numerical Variables

```
In [10]: numeric_features = train.select_dtypes(include=[np.number])
numeric_features.dtypes
```

```
Out[10]: User_ID          int64
Occupation          int64
Marital_Status      int64
Product_Category_1  int64
Product_Category_2  float64
Product_Category_3  float64
Purchase            int64
dtype: object
```


1.1.2.1. Distribution of the Occupation variable

```
In [11]: train.Occupation.value_counts()
```

```
Out[11]: 4      72308
0      69638
7      59133
1      47426
17     40043
20     33562
12     31179
14     27309
2      26588
16     25371
6      20355
3      17650
10     12930
5      12177
15     12165
11     11586
19      8461
13      7728
18      6622
9       6291
8       1546
Name: Occupation, dtype: int64
```

```
In [12]: sns.countplot(train.Occupation)
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfedfb14e0>
```




1.1.2.2. Distribution of the Marital_Status variable

```
In [13]: train.Marital_Status.value_counts()
```

```
Out[13]: 0    324731  
        1    225337  
        Name: Marital_Status, dtype: int64
```

```
In [14]: sns.countplot(train.Marital_Status)
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfedad3400>
```




1.1.2.3. Distribution of the Product_Category_1 variable

```
In [15]: train.Product_Category_1.value_counts()
```

```
Out[15]: 5      150933
         1      140378
         8      113925
        11      24287
         2      23864
         6      20466
         3      20213
         4      11753
        16       9828
        15       6290
        13       5549
        10       5125
        12       3947
         7       3721
        18       3125
        20       2550
        19       1603
        14       1523
        17        578
         9        410
        Name: Product_Category_1, dtype: int64
```

```
In [16]: sns.countplot(train.Product_Category_1)
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfedbd3400>
```




1.1.2.4. Distribution of the Product_Category_2 variable


```
In [17]: train.Product_Category_2.value_counts()
```

```
Out[17]: 8.0      64088
14.0     55108
2.0      49217
16.0     43255
15.0     37855
5.0      26235
4.0      25677
6.0      16466
11.0     14134
17.0     13320
13.0     10531
9.0       5693
12.0      5528
10.0      3043
3.0       2884
18.0      2770
7.0        626
Name: Product_Category_2, dtype: int64
```

```
In [18]: sns.countplot(train.Product_Category_2)
plt.xticks(rotation=90)
```

```
Out[18]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16]),
<a list of 17 Text xticklabel objects>)
```




1.1.2.5. Distribution of the Product_Category_3 variable

```
In [19]: train.Product_Category_3.value_counts()
```

```
Out[19]: 16.0    32636
15.0    28013
14.0    18428
17.0    16702
5.0     16658
8.0     12562
9.0     11579
12.0     9246
13.0     5459
6.0     4890
18.0     4629
4.0     1875
11.0     1805
10.0     1726
3.0       613
Name: Product_Category_3, dtype: int64
```

```
In [20]: sns.countplot(train.Product_Category_3)
plt.xticks(rotation=90)
```

```
Out[20]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]),
<a list of 15 Text xticklabel objects>)
```



1.1.2.6. Correlation between Numerical Predictors and Target variable


```
In [21]: corr = numeric_features.corr()

print (corr['Purchase'].sort_values(ascending=False)[:10], '\n')
print (corr['Purchase'].sort_values(ascending=False)[-10:])
```


```
Purchase          1.000000
Occupation        0.020833
User_ID           0.004716
Marital_Status   -0.000463
Product_Category_3 -0.022006
Product_Category_2 -0.209918
Product_Category_1 -0.343703
Name: Purchase, dtype: float64
```

```
Purchase          1.000000
Occupation        0.020833
User_ID           0.004716
Marital_Status   -0.000463
Product_Category_3 -0.022006
Product_Category_2 -0.209918
Product_Category_1 -0.343703
Name: Purchase, dtype: float64
```

```
In [22]: #correlation matrix
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corr, vmax=.9, square=True);
```



```
In [23]: #correlation matrix
f, ax = plt.subplots(figsize=(20, 9))
sns.heatmap(corr, vmax=.8,annot_kws={'size': 20}, annot=True);
```



```
In [24]: #Correlations Between Attributes  
#Pearson's Correlation  
#Coefficient, that assumes a normal distribution of the attributes involved"""  
  
s = corr.unstack()  
#s.sort_values(kind="quicksort")  
s
```

```

Out[24]: User_ID      User_ID      1.000000
          Occupation    -0.023971
          Marital_Status 0.020443
          Product_Category_1 0.003825
          Product_Category_2 0.001529
          Product_Category_3 0.003419
          Purchase      0.004716
Occupation User_ID      -0.023971
          Occupation    1.000000
          Marital_Status 0.024280
          Product_Category_1 -0.007618
          Product_Category_2 -0.000384
          Product_Category_3 0.013263
          Purchase      0.020833
Marital_Status User_ID      0.020443
          Occupation    0.024280
          Marital_Status 1.000000
          Product_Category_1 0.019888
          Product_Category_2 0.015138
          Product_Category_3 0.019473
          Purchase      -0.000463
Product_Category_1 User_ID      0.003825
          Occupation    -0.007618
          Marital_Status 0.019888
          Product_Category_1 1.000000
          Product_Category_2 0.540583
          Product_Category_3 0.229678
          Purchase      -0.343703
Product_Category_2 User_ID      0.001529
          Occupation    -0.000384
          Marital_Status 0.015138
          Product_Category_1 0.540583
          Product_Category_2 1.000000
          Product_Category_3 0.543649
          Purchase      -0.209918
Product_Category_3 User_ID      0.003419
          Occupation    0.013263
          Marital_Status 0.019473
          Product_Category_1 0.229678
          Product_Category_2 0.543649
          Product_Category_3 1.000000
          Purchase      -0.022006
Purchase User_ID      0.004716
          Occupation    0.020833
          Marital_Status -0.000463
          Product_Category_1 -0.343703
          Product_Category_2 -0.209918
          Product_Category_3 -0.022006
          Purchase      1.000000

dtype: float64


```

1.1.3. Categorical Variables

1.1.3.1. Distribution of the variable Gender

```
In [25]: sns.countplot(train.Gender)
```


```
Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfeef1bf60>
```



1.1.3.2. Distribution of the variable Age

```
In [26]: sns.countplot(train.Age)
```


```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfeef89c50>
```



1.1.3.3. Distribution of the variable City_Category

```
In [27]: sns.countplot(train.City_Category)
```


```
Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfeefdeb00>
```



1.1.3.4. Distribution of the variable Stay_In_Current_City_Years

```
In [28]: sns.countplot(train.Stay_In_Current_City_Years)
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfef016860>
```




1.2. Bivariate Distribution

1.2.1. Numerical Variables

1.2.1.1. Occupation and Purchase Analysis

```
In [29]: Occupation_pivot = \
train.pivot_table(index='Occupation', values="Purchase", aggfunc=np.mean)

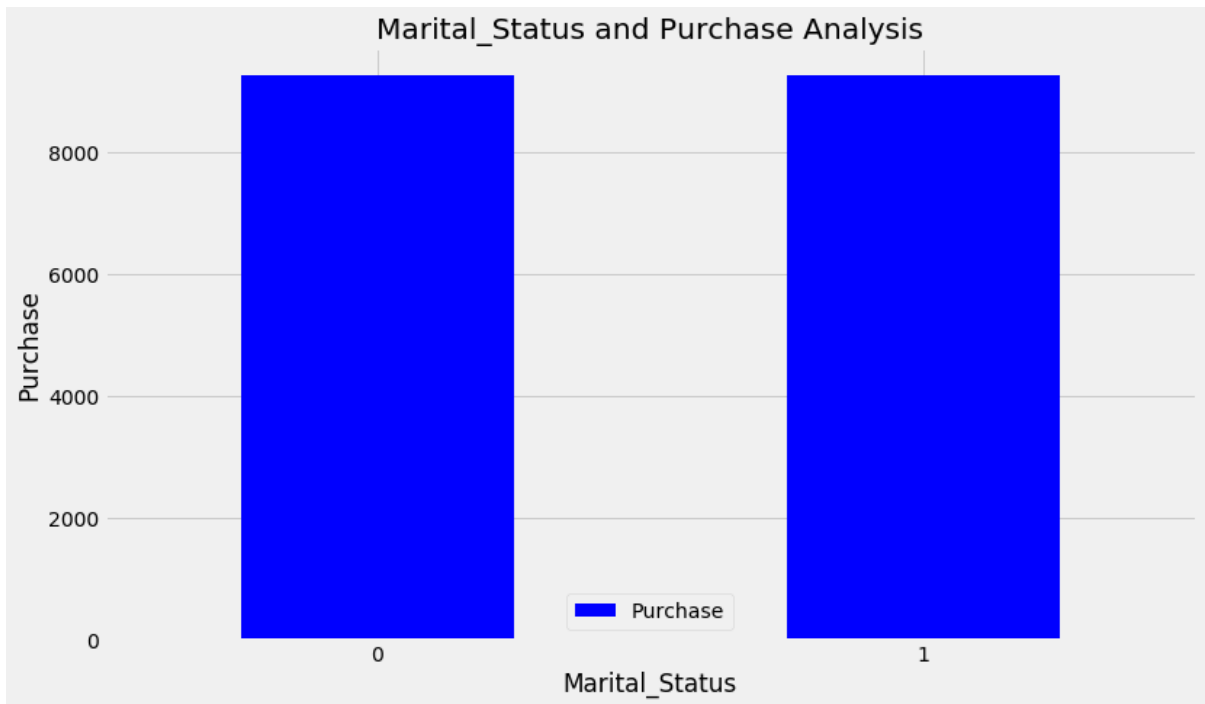
Occupation_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Occupation")
plt.ylabel("Purchase")
plt.title("Occupation and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



1.2.1.2. Marital_Status and Purchase Analysis

```
In [30]: Marital_Status_pivot = \
train.pivot_table(index='Marital_Status', values="Purchase", aggfunc=np.mean)


Marital_Status_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Marital_Status")
plt.ylabel("Purchase")
plt.title("Marital_Status and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



1.2.1.3. Product_Category_1 and Purchase Analysis


```
In [31]: Product_category_1_pivot = \
train.pivot_table(index='Product_Category_1', values="Purchase", aggfunc=np.me
an)

Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Product_Category_1")
plt.ylabel("Purchase")
plt.title("Product_Category_1 and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



```
In [32]: Product_category_1_pivot = \
train.pivot_table(index='Product_Category_1', values="Purchase", aggfunc=np.sum)


Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Product_Category_1")
plt.ylabel("Purchase")
plt.title("Product_Category_1 and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



1.2.1.4. Product_Category_2 and Purchase Analysis


```
In [33]: Product_category_1_pivot = \
train.pivot_table(index='Product_Category_2', values="Purchase", aggfunc=np.sum)

Product_category_1_pivot.plot(kind='bar', color='blue', figsize=(12,7))
plt.xlabel("Product_Category_2")
plt.ylabel("Purchase")
plt.title("Product_Category_2 and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



```
In [34]: Product_category_1_pivot = \
train.pivot_table(index='Product_Category_2', values="Purchase", aggfunc=np.me
an)


Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Product_Category_2")
plt.ylabel("Purchase")
plt.title("Product_Category_2 and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



1.2.1.4. Product_Category_3 and Purchase Analysis


```
In [35]: Product_category_1_pivot = \
train.pivot_table(index='Product_Category_3', values="Purchase", aggfunc=np.sum)

Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Product_Category_3")
plt.ylabel("Purchase")
plt.title("Product_Category_3 and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



```
In [36]: Product_category_1_pivot = \
train.pivot_table(index='Product_Category_3', values="Purchase", aggfunc=np.me
an)

Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Product_Category_3")
plt.ylabel("Purchase")
plt.title("Product_Category_3 and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```




1.2.2. Categorical Variables

1.2.2.1. Gender and Purchase Analysis


```
In [37]: Product_category_1_pivot = \
train.pivot_table(index='Gender', values="Purchase", aggfunc=np.mean)


Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Gender")
plt.ylabel("Purchase")
plt.title("Gender and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



1.2.2.2. Age and Purchase Analysis

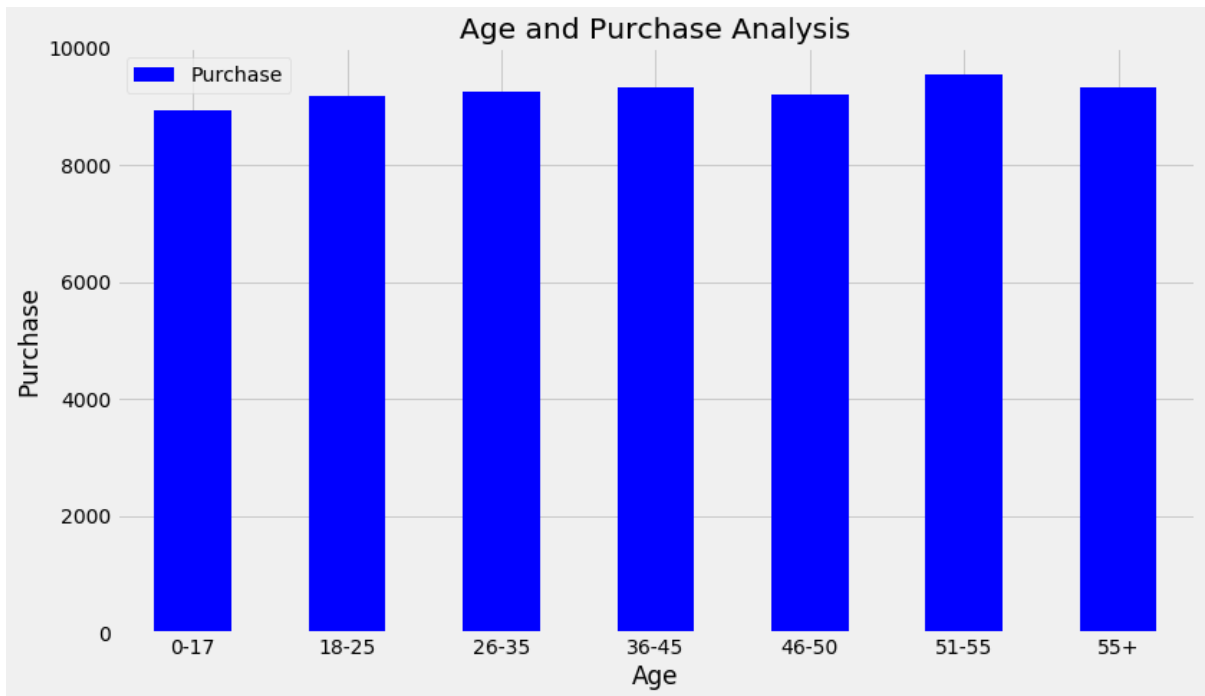
```
In [38]: Product_category_1_pivot = \
train.pivot_table(index='Age', values="Purchase", aggfunc=np.sum)

Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Age")
plt.ylabel("Purchase")
plt.title("Age and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



```
In [39]: Product_category_1_pivot = \
train.pivot_table(index='Age', values="Purchase", aggfunc=np.mean)


Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Age")
plt.ylabel("Purchase")
plt.title("Age and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



1.2.2.3. City_Category and Purchase Analysis

```
In [40]: Product_category_1_pivot = \
train.pivot_table(index='City_Category', values="Purchase", aggfunc=np.mean)


Product_category_1_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("City_Category")
plt.ylabel("Purchase")
plt.title("City_Category and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



1.2.2.4. Stay_in_Current_City_Years and Purchase Analysis

```
In [41]: Product_category_1_pivot = \
train.pivot_table(index='Stay_In_Current_City_Years', values="Purchase", aggfunc=np.mean)

Product_category_1_pivot.plot(kind='bar', color='blue', figsize=(12,7))
plt.xlabel("Stay_in_Current_City_Years")
plt.ylabel("Purchase")
plt.title("Stay_in_Current_City_Years and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



2. Data Preprocessing

```
In [42]: train.isnull().sum()
```

```
Out[42]: User_ID          0
Product_ID          0
Gender              0
Age                0
Occupation          0
City_Category       0
Stay_In_Current_City_Years  0
Marital_Status      0
Product_Category_1   0
Product_Category_2  173638
Product_Category_3  383247
Purchase            0
dtype: int64
```

```
In [43]: test.isnull().sum()
```

```
Out[43]: Unnamed: 0      0
         User_ID      0
         Product_ID   0
         Gender      0
         Age         0
         Occupation   0
         City_Category 0
         Stay_In_Current_City_Years 0
         Marital_Status 0
         Product_Category_1 0
         Product_Category_2 72344
         Product_Category_3 162562
         Comb        0
         dtype: int64
```

```
In [44]: print(train['Age'].value_counts())
         print(test['Age'].value_counts())
```

```
26-35    219587
36-45    110013
18-25     99660
46-50     45701
51-55     38501
55+       21504
0-17      15102
Name: Age, dtype: int64
26-35     93428
36-45     46711
18-25     42293
46-50     19577
51-55     16283
55+        9075
0-17       6232
Name: Age, dtype: int64
```

```
In [45]: train['Occupation'].value_counts() # 21
```

```
Out[45]: 4      72308
         0      69638
         7      59133
         1      47426
        17      40043
        20      33562
        12      31179
        14      27309
         2      26588
        16      25371
         6      20355
         3      17650
        10      12930
         5      12177
        15      12165
        11      11586
        19       8461
        13       7728
        18       6622
         9       6291
         8       1546
        Name: Occupation, dtype: int64
```

```
In [46]: train['City_Category'].value_counts()
```

```
Out[46]: B      231173
         C      171175
         A      147720
        Name: City_Category, dtype: int64
```

```
In [47]: train['Stay_In_Current_City_Years'].value_counts()
```

```
Out[47]: 1      193821
         2      101838
         3       95285
        4+       84726
         0       74398
        Name: Stay_In_Current_City_Years, dtype: int64
```

```
In [48]: train['Marital_Status'].value_counts()
```

```
Out[48]: 0      324731
         1      225337
        Name: Marital_Status, dtype: int64
```

```
In [49]: train['Product_Category_1'].value_counts() # 20
```

```
Out[49]: 5      150933
          1      140378
          8      113925
          11     24287
          2      23864
          6      20466
          3      20213
          4      11753
          16      9828
          15      6290
          13      5549
          10      5125
          12      3947
          7       3721
          18      3125
          20      2550
          19      1603
          14      1523
          17       578
          9       410
          Name: Product_Category_1, dtype: int64
```

```
In [50]: train['Product_Category_2'].value_counts() # 17
```

```
Out[50]: 8.0      64088
          14.0     55108
          2.0      49217
          16.0     43255
          15.0     37855
          5.0      26235
          4.0      25677
          6.0      16466
          11.0     14134
          17.0     13320
          13.0     10531
          9.0       5693
          12.0     5528
          10.0     3043
          3.0       2884
          18.0     2770
          7.0        626
          Name: Product_Category_2, dtype: int64
```



```
In [51]: train['Product_Category_3'].value_counts() # 15
```

```
Out[51]: 16.0    32636
        15.0    28013
        14.0    18428
        17.0    16702
        5.0     16658
        8.0     12562
        9.0     11579
        12.0     9246
        13.0     5459
        6.0     4890
        18.0     4629
        4.0     1875
        11.0     1805
        10.0     1726
        3.0       613
        Name: Product_Category_3, dtype: int64
```

```
In [52]: train.columns
```

```
Out[52]: Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation', 'City_Category',
               'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category_1',
               'Product_Category_2', 'Product_Category_3', 'Purchase'],
              dtype='object')
```

```
In [53]: test.columns
```

```
Out[53]: Index(['Unnamed: 0', 'User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation',
               'City_Category', 'Stay_In_Current_City_Years', 'Marital_Status',
               'Product_Category_1', 'Product_Category_2', 'Product_Category_3',
               'Comb'],
              dtype='object')
```

```
In [54]: train_join_test = train[['User_ID', 'Product_ID']]
train = train[['Gender', 'Age', 'Occupation', 'City_Category', 'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category_1', 'Purchase']]
train.head()
```

Out[54]:

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	F	0-17	10	A	2	0
1	F	0-17	10	A	2	0
2	F	0-17	10	A	2	0
3	F	0-17	10	A	2	0
4	M	55+	16	C	4+	0

```
In [55]: test = test[['Gender', 'Age', 'Occupation', 'City_Category', 'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category_1']]
test.head()
```

Out[55]:

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	M	46-50	7	B	2	0
1	M	26-35	17	C	0	0
2	F	36-45	1	B	4+	0
3	F	36-45	1	B	4+	0
4	F	26-35	1	C	1	0

3. Feature Engineering

```
In [56]: from sklearn.preprocessing import LabelEncoder
```

```
In [57]: le = LabelEncoder()
```

```
In [58]: train['Gender'] = le.fit_transform(train['Gender'])
        test['Gender'] = le.fit_transform(test['Gender'])
```

```
In [59]: train.head()
```

```
Out[59]:
```

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	0	0-17	10	A	2	0
1	0	0-17	10	A	2	0
2	0	0-17	10	A	2	0
3	0	0-17	10	A	2	0
4	1	55+	16	C	4+	0



```
In [60]: test.head()
```

```
Out[60]:
```

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	1	46-50	7	B	2	0
1	1	26-35	17	C	0	0
2	0	36-45	1	B	4+	0
3	0	36-45	1	B	4+	0
4	0	26-35	1	C	1	0



```
In [61]: train['Age'] = le.fit_transform(train['Age'])
        test['Age'] = le.fit_transform(test['Age'])
```

```
In [62]: train['City_Category'] = le.fit_transform(train['City_Category'])
```

```
In [63]: test['City_Category'] = le.fit_transform(test['City_Category'])
```

In [64]: train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 8 columns):
Gender                550068 non-null int32
Age                  550068 non-null int32
Occupation            550068 non-null int64
City_Category         550068 non-null int32
Stay_In_Current_City_Years  550068 non-null object
Marital_Status        550068 non-null int64
Product_Category_1    550068 non-null int64
Purchase              550068 non-null int64
dtypes: int32(3), int64(4), object(1)
memory usage: 27.3+ MB
```

In [65]: test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 233599 entries, 0 to 233598
Data columns (total 7 columns):
Gender                233599 non-null int32
Age                  233599 non-null int32
Occupation            233599 non-null int64
City_Category         233599 non-null int32
Stay_In_Current_City_Years  233599 non-null object
Marital_Status        233599 non-null int64
Product_Category_1    233599 non-null int64
dtypes: int32(3), int64(3), object(1)
memory usage: 9.8+ MB
```

```
In [66]: mask = train['Stay_In_Current_City_Years'] == '4+'
train.loc[mask, 'Stay_In_Current_City_Years'] = 4
```

```
In [67]: mask1 = test['Stay_In_Current_City_Years'] == '4+'
test.loc[mask1, 'Stay_In_Current_City_Years'] = 4
```

In [68]: train.head()


Out[68]:

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	0	0	10	0	2	0
1	0	0	10	0	2	0
2	0	0	10	0	2	0
3	0	0	10	0	2	0
4	1	6	16	2	4	0

In [69]: `test.head()`

Out[69]:

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	1	4	7	1	2	0
1	1	2	17	2	0	0
2	0	3	1	1	4	0
3	0	3	1	1	4	0
4	0	2	1	2	1	0



In [70]: `train['Stay_In_Current_City_Years'] = train['Stay_In_Current_City_Years'].astype('int')`
`test['Stay_In_Current_City_Years'] = test['Stay_In_Current_City_Years'].astype('int')`

In [71]: `train.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 8 columns):
Gender                550068 non-null int32
Age                  550068 non-null int32
Occupation            550068 non-null int64
City_Category         550068 non-null int32
Stay_In_Current_City_Years  550068 non-null int32
Marital_Status        550068 non-null int64
Product_Category_1     550068 non-null int64
Purchase              550068 non-null int64
dtypes: int32(4), int64(4)
memory usage: 25.2 MB
```

In [72]: `test.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 233599 entries, 0 to 233598
Data columns (total 7 columns):
Gender                233599 non-null int32
Age                  233599 non-null int32
Occupation            233599 non-null int64
City_Category         233599 non-null int32
Stay_In_Current_City_Years  233599 non-null int32
Marital_Status        233599 non-null int64
Product_Category_1     233599 non-null int64
dtypes: int32(4), int64(3)
memory usage: 8.9 MB
```

4. Model

4.1 Multiple Linear Regression

```
In [75]: X_train = train.drop('Purchase',axis=1)
y_train = train['Purchase']
X_test = test
```

```
In [76]: from sklearn.linear_model import LinearRegression
```

```
In [77]: reg = LinearRegression()
```

```
In [78]: reg.fit(X_train, y_train)
```

```
Out[78]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

```
In [89]: prediction1 = reg.predict(X_test)
```

```
In [90]: prediction1
```

```
Out[90]: array([11492.07110442, 10785.26962272, 9091.44149746, ...,
10815.25902604, 7372.75994218, 9638.41030116])
```

```
In [91]: pred = pd.DataFrame(prediction1, columns = ['Purchase'])
pred.head()
```

```
Out[91]:
```

	Purchase
0	11492.071104
1	10785.269623
2	9091.441497
3	9528.730516
4	9739.615697

```
In [94]: final = pd.concat([train_join_test,pred],axis=1)
final.head()
```

```
Out[94]:
```

	User_ID	Product_ID	Purchase
0	1000001	P00069042	11492.071104
1	1000001	P00248942	10785.269623
2	1000001	P00087842	9091.441497
3	1000001	P00085442	9528.730516
4	1000002	P00285442	9739.615697

```
In [95]: final.to_csv('finalpurchase-MLR.csv', index = False) # RMSE: 4713.9227352998 , Rank: 1344
```

4.2 Decision Tree Regression

```
In [96]: from sklearn.tree import DecisionTreeRegressor
regressor = DecisionTreeRegressor(random_state = 0)
regressor.fit(X_train, y_train)
```

```
Out[96]: DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                presort=False, random_state=0, splitter='best')
```

```
In [97]: prediction3 = regressor.predict(X_test)
pred = pd.DataFrame(prediction3, columns = ['Purchase'])
pred.head()
```

Out[97]:

	Purchase
0	16599.531250
1	10205.857143
2	6984.247191
3	2372.000000
4	2184.375000

```
In [99]: final = pd.concat([train_join_test, pred], axis=1)
final.head()
```

Out[99]:

	User_ID	Product_ID	Purchase
0	1000001	P00069042	16599.531250
1	1000001	P00248942	10205.857143
2	1000001	P00087842	6984.247191
3	1000001	P00085442	2372.000000
4	1000002	P00285442	2184.375000

```
In [100]: final.to_csv('finalpurchase-DTR.csv', index = False) # RMSE: , Rank:
```

4.3 Random Forest Regression

```
In [101]: from sklearn.ensemble import RandomForestRegressor
reg = RandomForestRegressor(n_estimators = 10, random_state = 0)
reg.fit(X_train, y_train)
```

```
Out[101]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)
```

```
In [102]: prediction4 = regressor.predict(X_test)
pred = pd.DataFrame(prediction4, columns = ['Purchase'])
pred.head()
```

```
Out[102]:
```

	Purchase
0	16599.531250
1	10205.857143
2	6984.247191
3	2372.000000
4	2184.375000

```
In [104]: final = pd.concat([train_join_test, pred], axis=1)
final.head()
```

```
Out[104]:
```

	User_ID	Product_ID	Purchase
0	1000001	P00069042	16599.531250
1	1000001	P00248942	10205.857143
2	1000001	P00087842	6984.247191
3	1000001	P00085442	2372.000000
4	1000002	P00285442	2184.375000

```
In [105]: final.to_csv('finalpurchase-RFR.csv', index = False) # RMSE: 3119.9813685216
, Rank: 1143
```

ANN

```
In [ ]: #import keras
from keras.models import Sequential
from keras.layers import Dense
```



```
In [ ]: ann = Sequential()

In [ ]: ann.add(Dense(6, init = 'uniform', activation = 'relu', input_dim = 7))

In [ ]: ann.add(Dense(6, init = 'uniform', activation = 'relu'))

In [ ]: ann.add(Dense(1, activation='linear'))

In [ ]: ann.compile(loss='mse', optimizer='adam', metrics=['mse','mae'])

In [ ]: ann.fit(X_train, y_train, batch_size = 10, nb_epoch = 10)

In [ ]: prediction5 = ann.predict(X_test)
        pred = pd.DataFrame(prediction5, columns = ['Purchase'])
        pred.head()

In [ ]: final = pd.concat([df_join_test,pred],axis=1)
        final.head()

In [ ]: final.to_csv('finalpurchase-ANN.csv', index = False) # RMSE:    , Rank:
```

Evaluation Results

Multiple Linear Regression- RMSE: 4713.9227352998 , Rank: 1344

Random Forest Regression- RMSE: 3119.9813685216 , Rank: 1143